Keyword Search over Data Service Integration for Accurate Results



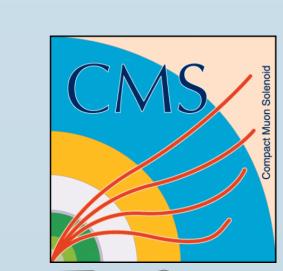


Vidmantas Zemleris

Vilnius University, Lithuania vidmantas.zemleris@cern.ch for the benefit of CMS Collaboration

Valentin Kuznetsov

Cornell University, USA vkuznet@gmail.com





Summary

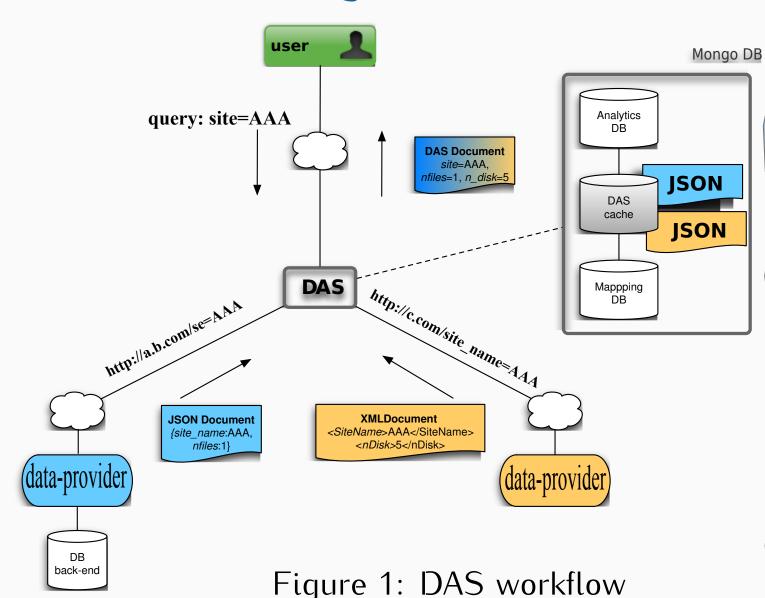
Virtual data integration aims at providing a coherent interface for querying heterogeneous data sources (e.g. web services, proprietary systems) with minimum upfront effort in integration. Data is usually accessed through a structured queries, such as SQL, requiring to learn the language and to get acquainted with data organization, which may pose problems even to proficient users.

We present a keyword search system, which proposes a ranked list of structured queries along with their explanations. It operates mainly on the metadata, such as the constraints on inputs accepted by services. It was developed as an integral part of the CMS data discovery service and is currently available as open source.

Context: a system for Virtual Data Integration

"CMS Data Aggregation System" (DAS):

- accepts simple structured queries
- integrates heterogeneous services
- → parse the query
- → contact services
- → eliminate inconsistencies in the responses: * entity naming; data formats (XML, JSON)
- → combine the responses
- requires only minimal service mappings
- → no predefined schema
- → minimal effort in defining services



Queries must specify: entity to be retrieved and filtering criteria. Optionally, the results can be further filtered, sorted or aggregated dataset=*RelVal* | grep dataset.nevents >1000 | avg(dataset.size), median(dataset.size)

entity requested from services

conditions as service *inputs* filters and projections on service *outputs*

aggregators

still, it is overwhelming for users to:

- learn the query language
- remember how exactly the data is structured and named
- Could keyword queries solve this?

Interpreting Keyword Queries: Problem definition

Input: query, KWQ= $(kw_1, kw_2, ..., kw_n)$

ambigous; nearby keywords are often related

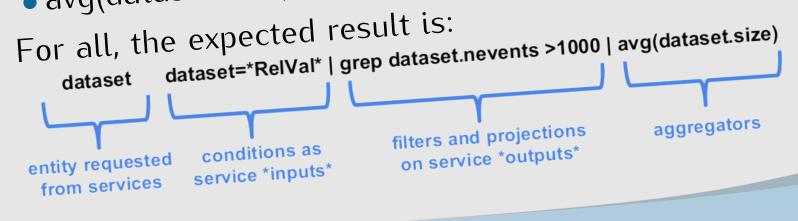
Task: translate it into structured query

(made of tag_i - schema and value terms, unknown, operators)

GIVEN: metadata only:

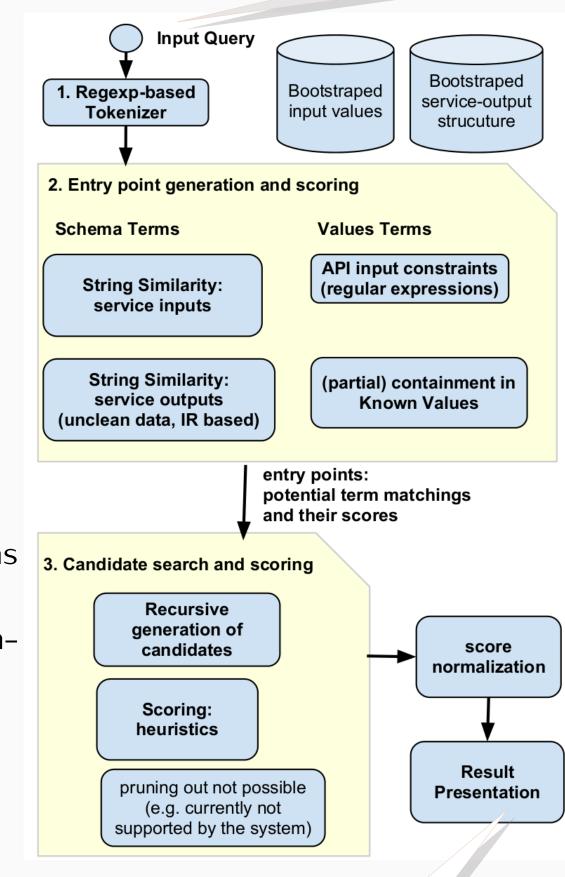
- names of entities and their attributes
- → service inputs or their output fields
- possible values (only for some inputs)
- constraints on data-service inputs:
- → mandatory inputs
- → regular expressions on values

Example. Consider this query: average size of RelVal datasets with its number of events > 1000 average RelVal dataset size nevents>1000 avg(dataset size) RelVal "number of events">1000



Keyword search overview

- tokenizer:
- \rightarrow clean up the query
- → identify patterns
- identify and score "entry points" with
- → string matching [for entity names]
- → IR (IDF-based) [output fieldnames]
- → list of known values
- → regular expressions on allowed values
- combine *entry points*
- \rightarrow consider various entry point permutations 3. Candidate search and scoring (keyword labelings)
- → promote ones respecting keyword dependencies or other heuristics
- → interpret as structured queries



Did you mean any of the queries below?

Filter by entity: dataset, file, summary, block, lumi, any

ts>100



Challenges

- keyword queries are ambiguous → return ranked list of structured query suggestions
- no direct access to the data querying services is "expensive"
- → rely on metadata
- * bootstrap list of allowed values (available only for some fields)
- * rely on *regexps* with lower confidence (can result in false positives)
- no predefined schema
- → bootstrap list of fields in service results through queries
- \rightarrow some field names are unclean \rightarrow use IDF (as they come directly from JSON/XML responses)

The ranker

Use exhaustive search:

- allows easily finding optimal solutions, vs. complex methods that'd require post-pruning
- early pruning filter out many "invalid" candidates e.g. not yet supported by services
- our schema is quite small
- \rightarrow cython-based implementation is already quite fast (bound by MongoDB/IR engine to retrieve entry points)

Scoring function

$$final\ score = \sum_{i=1}^{|KWQ|} \left(\log \left(score_{tag_i|kw_i} \right) + \sum_{h_j \in H} h_j(tag_i|kw_i; tag_{i-1,..,1}) \right)$$

 $score_{tag_i|kw_i}$ - likelihood of kw_i to be tag_i (from entry points step)

 $h_i(tag_i|kw_i;tag_{i-1,..,1})$ - the score boost returned by heuristic h_i given the tag(s) nearby.

Our finding: summing log-likelihoods is better than plain scores (cf. Keymantic) $ascore(tag_i = unknown|kw_i)$ currently uses predefined P(unknown).

Related works

- Keymantic (the closest work)
- 1. score keyword mappings individually (entry points)
- 2. solve "weighted bipartite assignment" $(kw_i \rightarrow tag_i)$ with contextualizations:
- → maximize total sum of weights
- → uses heuristics to account for keyword interdependencies (contextualization)
 - * e.g. <table_name> <attribute>; <attribute> <its value>;
- * solves it approximately with Munkres algorithm modified to consider contextualizations:
- · contextualize modify weights of $kw_i \rightarrow tag_i$, if tag_i is "related" to earlier sub-assignments \cdot to get multiple results, repeat recursively forcing/preventing certain sub-assignments
- 3. interpret generated mappings as SQL queries
- KEYRY uses HMM (Hidden Markov Model) to label keywords as schema terms
- → HMM's initial parameters can be estimated from similar heuristics as above
- → later machine learning can be used (if logs available)

Autocompletion to ease typing queries (prototype)

ataset=/AlCaLumiPixels/T0TEST_mixed_tl_processed_524pl_BUNNIES-Tie Q dataset=/DoubleElectron/T0TEST 532p1Run2012C BUNNIES-DiTau-Tier1Pr Site=Baylor-Tier3

tter is a data tier used by DBS system to identify dataset meaning, there are different pre-defined data-tiers, such as GEN, SIM, RECO, ALCORECO, RAW and mixed one, e.g. GEN-SIM-» dataset tier=*GEN* Relationships (i.e. can retrieve these below given tier):

Future work

- improve autocompletion prototype
- improve the ranker
- generic ways to improve services' performance, e.g. materialized views with incremental refresh

Open problems & ideas

TOP-K (SEMI-)OPTIMAL SOLUTIONS WITH CONTEXTUALIZATION?

- can we adapt Murty's/Munkres's algs. for top-k list of assignments with contextualizations?
- → this would at least guarantee optimal top-k for with **some** contextualization
- → out of scope, ask for handouts/chat

PROBLEMATICS OF THE HMM APPROACH:

- what is modelled is not necessarily same as seen by user
- \rightarrow models $kw_i \rightarrow tag_j$, while user sees structured queries
- → therefore, hard to automatically collect training data

